

Multi-Layer Neural Networks

[RN2] Sec 20.5

[RN3] Sec 20.5

CS 486/686


University of Waterloo

Lecture 20: July 9, 2015

Multilayer Feed-forward Neural Networks

- Perceptron can only represent (soft) linear separators
 - Because single layer
- With multiple layers, what fns can be represented?
 - Virtually any function!


Multilayer Networks



$$a_i = g\left(\sum_j W_{ji} g\left(\sum_k W_{kj} a_k\right)\right)$$

Multilayer Networks


- Adding two sigmoid units with parallel but opposite "cliffs" produces a ridge



Multilayer Networks

- Adding two intersecting ridges (and thresholding) produces a bump

Network output



Multilayer Networks

- By tiling bumps of various heights together, we can approximate any function
- **Theorem:** Neural networks with at least one hidden layer of sufficiently many sigmoid units can approximate any function arbitrarily closely.

Common Activation Functions

- Threshold: $h(x) = \begin{cases} 1 & x \geq 0 \\ -1 & x < 0 \end{cases}$
- Sigmoid: $h(x) = \sigma(x) = \frac{1}{1+e^{-x}}$
- Gaussian: $h(x) = e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$
- Hyperbolic tangent: $h(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- Identity: $h(x) = x$

Weight Training

- Parameters: $\langle W^{(1)}, W^{(2)}, \dots \rangle$
- Objectives:
 - Error minimization
 - Backpropagation (aka “backprop”)
 - Maximum likelihood
 - Maximum a posteriori
 - Bayesian learning

Least squared error

- Error function

$$E(\mathbf{W}) = \frac{1}{2} \sum_n E_n(\mathbf{W})^2 = \frac{1}{2} \sum_n ||f(\mathbf{x}_n, \mathbf{W}) - y_n||_2^2$$

where \mathbf{x}_n is the input of the n^{th} example

y_n is the label of the n^{th} example

$f(\mathbf{x}_n, \mathbf{W})$ is the output of the neural net

Sequential Gradient Descent

- For each example (x_n, y_n) adjust the weights as follows:

$$W_{ji} \leftarrow W_{ji} - \alpha \frac{\partial E_n}{\partial W_{ji}}$$

- How can we compute the gradient efficiently given an arbitrary network structure?
- Answer: **backpropagation algorithm**

Backpropagation

- Back-Prop-Learning(examples, network)
 - Repeat
 - For each example e do
 - Compute output a of each node in **forward** pass
 - » Input nodes: $a_j \leftarrow x_j[e]$
 - » Other nodes: $in_i \leftarrow \sum_j W_{ji} a_j$ and $a_i \leftarrow g(in_i)$
 - Compute modified error Δ of each node in **backward** pass ($l = L$ to 1)
 - » Output nodes: $\Delta_i \leftarrow g'(in_i) (y_i[e] - a_i)$
 - » For each node j in layer l : $\Delta_j \leftarrow g'(in_j) \sum_i W_{ji} \Delta_i$
 - » For each node i in layer $l + 1$: $W_{ji} \leftarrow W_{ji} + \alpha a_j \Delta_i$
 - Until some stopping criteria satisfied
 - Return learnt network

Forward phase

- Propagate inputs forward to compute the output of each unit
- Output a_i at unit i :

$$a_i = g(in_i) \quad \text{where} \quad in_i = \sum_j W_{ji} a_j$$

Backward phase

- Use chain rule to recursively compute gradient

- For each weight W_{ji} : $\frac{\partial E_n}{\partial W_{ji}} = \frac{\partial E_n}{\partial in_i} \frac{\partial in_i}{\partial W_{ji}} = \Delta_i a_j$

- Let $\Delta_i \equiv \frac{\partial E_n}{\partial in_i}$ then

$$\Delta_i = \begin{cases} g'(in_i)(y_i - a_i) & \text{base case: } i \text{ is an output unit} \\ g'(in_i) \sum_j W_{ji} \Delta_j & \text{recursion: } i \text{ is a hidden unit} \end{cases}$$

- Since $in_i = \sum_j W_{ji} a_j$ then $\frac{\partial in_i}{\partial W_{ji}} = a_j$

Simple Example

- Consider a network with two layers:
 - Hidden nodes: $g(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
 - Tip: $\tanh'(x) = 1 - \tanh^2(x)$
 - Output node: $g(x) = x$
- Objective: squared error


Simple Example

- Forward propagation:
 - Hidden units: $in_j =$ $a_j =$
 - Output units: $in_i =$ $a_i =$
- Backward propagation:
 - Output units: $\Delta_i =$
 - Hidden units: $\Delta_j =$
- Gradients:
 - Hidden layers: $\frac{\partial E_n}{\partial W_{kj}} =$
 - Output layer: $\frac{\partial E_n}{\partial W_{ji}} =$


Non-linear regression examples

- Two layer network:
 - 3 tanh hidden units and 1 identity output unit


$$y = x^2$$




$$y = \sin x$$



$$y = |x|$$



$$y = \int_{-\infty}^x \delta(t) dt$$



Analysis

- Efficiency:
 - Fast gradient computation: linear in number of weights
- Convergence:
 - Slow convergence (linear rate)
 - May get trapped in local optima
- Prone to overfitting
 - Solutions: early stopping, regularization (add $\|w\|_2^2$ penalty term to objective)

Neural Net Applications

- Neural nets can approximate any function, hence 1000's of applications
 - Speech recognition
 - Character recognition
 - Paint-quality inspection
 - Vision-based autonomous driving
 - Etc.